

Training Day 2 Report

Date: 2 July 2025

By: Sahil Stathia Sharma

URN: 2302658

CRN: 2315255

Overview

On the second day of the training session, we focused on Large Language Models (LLMs), a core area of artificial intelligence. These models are capable of understanding and generating human-like text based on vast amounts of training data. The session gave us insight into how LLMs work, their components, the training process, and their impact across different sectors like education, healthcare, customer support, and creative industries. Real-world examples such as ChatGPT, Claude, and Gemini were discussed to help us connect theory to practice. The session also introduced ethical concerns and limitations associated with AI technologies.

We Explored

Definition and real-world analogy of LLMs

Key terms like Token, Parameter, Prompt, Fine-tuning, Inference

Working process of LLMs through tokenization and transformer models

Evolution of LLMs from GPT-1 to GPT-4o and Gemini

Three main stages of training: Pretraining, Fine-tuning, and RLHF

Prompt Engineering types: Zero-shot, Few-shot, Chain-of-thought

Real-world applications in education, healthcare, customer support, law, and content creation

Limitations like hallucinations, biases, and lack of actual understanding

Ethical concerns and the future of LLMs in multimodal AI, autonomous agents, and mobile deployment

1. What is a Large Language Model?

A Large Language Model (LLM) is a form of artificial intelligence trained to understand and generate text in natural human language. It analyzes enormous datasets—ranging from books to internet conversations—to learn grammar, context, and semantics.

Examples:

ChatGPT (OpenAI)

Google Gemini

Claude (Anthropic)
LLaMA (Meta)

Analogy: LLM is like a human who has read billions of pages and can answer or generate text on nearly any topic.

2. Key Terms

Token: Smallest unit of text (word or word part)

Parameter: Model's internal settings used to make predictions

Prompt: The input or question given to the model

Fine-Tuning: Further training of a model on specific data

Inference: The model's generated output/response

Analogy: Tokens are like ingredients, parameters are recipe rules, and inference is the final dish.

3. Evolution of Language Models

Year	Model	Parameters	Creator
2018	BERT	110M	Google
2019	GPT-1	117M	OpenAI
2023	GPT-3	175B	OpenAI
2024	LLaMA, Claude	70B–100B	Meta, Anthropic
2024	GPT-4.0, Gemini 1.5	~200B	OpenAI, Google

4. How Do LLMs Work?

1. User Prompt: e.g., "What is the capital of France?"
2. Tokenization: ["What", "is", "the", "capital", "of", "France", "?"]
3. Processing: The model uses transformer layers to analyze tokens.
4. Prediction: The next token is predicted based on context → "Paris"
5. Output: A complete, relevant response is generated.

> Insight: LLMs don't "understand" like humans—they predict based on patterns from training data.

5. Training LLMs – Behind the Scenes

1. Pretraining: Learning general language patterns from large datasets
2. Fine-Tuning: Adapting the model to specific domains (e.g., medicine)
3. RLHF (Reinforcement Learning with Human Feedback): Human trainers help the model learn better responses

Analogy: Like teaching a parrot basic words, then coaching it to form meaningful sentences.

6. Applications of LLMs

Chatbots: 24/7 customer support

Education: AI-based tutoring

Healthcare: Medical note summarization

Legal: Drafting legal documents

Creative Writing: Poems, blogs, code, etc.

Demo Prompt Example:

"Write a funny poem about a tired student."

7. Prompt Engineering Techniques

1. Zero-shot Prompting

Asking without any prior examples

Example: "Translate 'Hello' to Spanish."

2. Few-shot Prompting

Providing input-output examples

Example: English: Hello → Spanish: Hola

3. Chain-of-Thought Prompting

Asking the model to explain step-by-step

Example: "Let's solve this problem step-by-step."

Tip: Be specific, clear, and role-based.

Example: "Act like a historian and explain World War I."

8. Limitations of LLMs

Hallucinations: May generate false or misleading information

Biases: Can reflect social or political biases from training data

No Actual Understanding: Predicts, doesn't "think"

Context Limitations: Memory is limited to a few thousand tokens

> Note: Always verify important information when using LLMs.

9. Ethical Concerns

- Spread of misinformation
- Violation of data privacy
- Risk of misuse (deepfakes, impersonation)
- Lack of accountability

Case Study: Microsoft's Tay chatbot turned offensive due to biased user input within hours of launch.

10. Future of LLMs

Multimodal Models: Understand and generate across text, image, and audio

Autonomous Agents: AI systems that perform tasks independently (e.g., Devin)

Lightweight Models: Smaller, faster models for public use

On-device LLMs: Local AI on phones and computers (e.g., Apple, Android devices)